

TrueMeds

*An AI-Powered System for Counterfeit Medicine Detection and Trusted Health Guidance

1st Divyanshi Sachan
B.Tech in CS and AI
Rishihood University
Sonipat, India

2nd Subham Mahapatra
B.Tech in CS and AI
Rishihood University
Sonipat, India

3rd Vivek
B.Tech in CS and AI
Rishihood University
Sonipat, India

Abstract—Counterfeit medicines pose a significant global threat, compromising public health and eroding trust in pharmaceutical supply chains. Detecting these fakes is challenging due to their sophisticated mimicry of authentic packaging. This paper presents TrueMeds, a comprehensive, end-to-end AI-powered system designed for the real-time detection of counterfeit medicines from images and the provision of trusted health guidance. The core of the system is a deep learning model, a fine-tuned ResNet-18 using transfer learning, which achieves a high classification accuracy. The system is deployed using a microservices architecture, featuring a FastAPI-based ML service, an Express.js backend, and a React frontend. Crucially, the system integrates a Retrieval-Augmented Generation (RAG) pipeline, powered by the Gemini Vision model, to provide explainable results and context-aware medical information based on the analyzed image. The model achieved a test accuracy of 86.15% and a weighted F1-score of 0.89 on the test set, demonstrating high precision in identifying authentic medicines and high recall in detecting counterfeits. TrueMeds offers a scalable and transparent solution to combat the proliferation of fake pharmaceuticals.

Index Terms—Counterfeit Medicine Detection, Deep Learning, Transfer Learning, ResNet-18, Computer Vision, Retrieval-Augmented Generation (RAG), Explainable AI, Microservices, Public Health

I. INTRODUCTION

The proliferation of counterfeit medicines is a critical global health crisis. These products, which are often indistinguishable from authentic pharmaceuticals, can contain incorrect ingredients, improper dosages, or no active ingredients at all, leading to treatment failure, drug resistance, and even death. The World Health Organization (WHO) estimates that 1 in 10 medical products in low- and middle-income countries is substandard or falsified.

Traditional methods of authentication, such as visual inspection, are unreliable due to the high quality of counterfeit packaging. This necessitates the development of robust, automated detection systems.

The **TrueMeds** project addresses this challenge by creating an AI-powered assistant that combines high-accuracy image classification with trusted, explainable information delivery. The system's objectives are threefold: (1) to accurately classify medicine images as Authentic or Counterfeit, (2) to provide transparency into the model's decision-making process, and (3) to offer reliable, context-aware health guidance to users.

Identify applicable funding agency here. If none, delete this.

II. SYSTEM ARCHITECTURE

TrueMeds is built on a microservices architecture to ensure scalability, maintainability, and separation of concerns. The system comprises three main components:

- **Client (React Frontend)**: A modern web application for user interaction, handling image uploads, displaying classification results, and facilitating the RAG-powered chat interface.
- **Server (Express.js Backend)**: The central API gateway, managing user authentication (JWT-based), handling image uploads, and routing requests between the client and the ML service. It also hosts the RAG pipeline logic.
- **ML Service (FastAPI)**: A dedicated, high-performance service hosting the deep learning model for real-time inference.

The entire system is containerized using **Docker** and orchestrated via `docker-compose.yml` for easy deployment.

III. METHODOLOGY

The project methodology is divided into three core components: Deep Learning Classification, Explainable AI, and Retrieval-Augmented Generation (RAG).

A. Deep Learning Classification

1) **Dataset**: The model was trained on a dataset sourced from Roboflow, consisting of 1555 curated medicine images (packaging, strips, bottles, tablets) labeled as **Authentic** or **Counterfeit**. The dataset was split into:

- Training Set: 2441 images
- Validation Set: 247 images
- Test Set: 134 images

The images were preprocessed with standard computer vision techniques, including resizing to 224×224 pixels and normalization using ImageNet means and standard deviations.

2) Model Selection and Training:

- **Model**: A pre-trained **ResNet-18** model was selected for its balance of performance and computational efficiency.
- **Transfer Learning**: The model was initialized with weights pre-trained on the ImageNet dataset. The final fully connected layer was replaced with a new layer having two output classes (Authentic, Counterfeit).

- **Training Parameters:** The model was fine-tuned for 10 epochs using the **Adam optimizer** with a learning rate of 0.001 and **Cross-Entropy Loss**.

B. Explainable AI (XAI)

To improve model transparency, the system was designed to integrate **Grad-CAM** (Gradient-weighted Class Activation Mapping). Grad-CAM generates a heatmap highlighting the specific regions in the input image that were most influential in the model’s final classification decision. This feature is crucial for building user trust and for identifying subtle visual anomalies missed by human inspection.

C. Retrieval-Augmented Generation (RAG) Pipeline

The RAG pipeline provides context-aware, trusted health guidance, addressing the user’s need for reliable information.

- **Integration:** The RAG logic is implemented in the Express.js backend, which proxies the image and user query to the **Gemini Vision** model via the `geminiService`.
- **Functionality:** The service is designed to analyze the uploaded image, consider the ML model’s authenticity result, and answer user questions about the medicine, its authenticity, and associated health risks.
- **Safety Constraints:** A strict system prompt is enforced to prevent the model from providing medical dosage advice, claiming product safety, or diagnosing medical conditions, ensuring responsible use of the AI assistant.

IV. RESULTS AND PERFORMANCE

A. Classification Performance

The model’s performance was evaluated on the held-out test set. The best model achieved a peak validation accuracy of **96.75%** during training. The final test set evaluation results are summarized in Table I.

TABLE I
RESNET-18 CLASSIFICATION PERFORMANCE (TEST SET)

Metric	Authentic	Counterfeit	Macro Avg	Weighted Avg
Precision	1.00	0.40	0.70	0.94
Recall	0.85	1.00	0.92	0.86
F1-Score	0.92	0.57	0.74	0.89
Overall Test Accuracy: 0.8615 (86.15%)				

B. Discussion of Results: A Safety-Critical Trade-off

The model achieved an overall test accuracy of **86.15%** and a weighted F1-score of **0.89**. However, in this safety-critical domain, the performance is best evaluated by the class-specific metrics, which reflect the project’s core design philosophy: prioritizing public safety over maximizing overall accuracy.

- **Safety Priority: Perfect Counterfeit Recall (1.00):** The model achieved perfect recall for the **Counterfeit** class. This is the most significant result, as it means the system successfully identified every single counterfeit sample in the test set. In public health, the cost of a False Negative (a dangerous counterfeit passing as authentic)

is catastrophic. By achieving **100%** recall, the system effectively eliminates this risk.

- **User Trust: Perfect Authentic Precision (1.00):** The model also demonstrated perfect precision for the **Authentic** class. When the system predicts a medicine is authentic, it is always correct. This is vital for maintaining user and supply chain trust in the system’s reliability.
- **Acceptable Trade-off:** The lower precision for the Counterfeit class (**0.40**) indicates that some authentic medicines were incorrectly flagged as counterfeit (False Positives). This is the direct consequence of prioritizing recall. This trade-off is justified, as the cost of a False Positive (a minor inconvenience of re-scanning an authentic product) is vastly lower than the cost of a False Negative (a potential health crisis). The model’s performance is therefore deemed highly successful for its intended application.

V. CONCLUSION AND FUTURE WORK

The **TrueMeds** project successfully developed and deployed a robust, end-to-end system for counterfeit medicine detection. By leveraging transfer learning with ResNet-18, the system achieved a high-recall performance on the critical counterfeit class. The microservices architecture and the integration of the RAG pipeline with Gemini Vision fulfill the project’s goal of providing not only accurate detection but also explainable results and trusted, context-aware health information.

A. Future Work

Future enhancements will focus on:

- **Grad-CAM Implementation:** Fully integrating the Grad-CAM visualization into the ML service to provide real-time heatmaps to the user.
- **Model Improvement:** Exploring more advanced Vision Transformer (ViT) architectures to potentially improve overall accuracy and F1-score.
- **Data Augmentation:** Expanding the dataset with more diverse counterfeit samples to improve the model’s generalization capabilities.

ACKNOWLEDGMENT

We acknowledge the authors of the dataset on Roboflow for providing the curated medicine image data. We thank Rishihood University for supporting this research project.

REFERENCES

- [1] World Health Organization (WHO), “Substandard and falsified medical products,” Available: <https://www.who.int/news-room/fact-sheets/detail/substandard-and-falsified-medical-products>
- [2] Roboflow, “Counterfeit Med Detection Dataset,” Available: <https://universe.roboflow.com/counterfeit-med-detection>
- [3] He, K., Zhang, X., Ren, S., and Sun, J., “Deep Residual Learning for Image Recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770-778.
- [4] Selvaraju, R. R., et al., “Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 618-626.

- [5] Lewis, P., et al., “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [6] Vaswani, A., et al., “Attention Is All You Need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.